

# Auto-Encoding Variational Bayes

Diederik P. Kingma and Max Welling

Ashish Bora

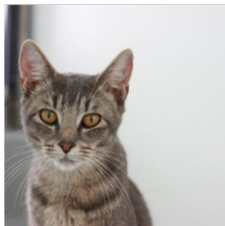
Oct 2016

# Outline

- 1 Quick intro to Bayesian Learning
- 2 Latent Variable Models
- 3 Tasks
- 4 Known Approaches
- 5 Variational Inference
- 6 Parametrization of distributions
- 7 Putting it together
- 8 Experiments and Results
- 9 Demo

- Observations  $x$ , Parameters/Unobserved variables  $z$   
Example

$x =$



,  $z = \text{cat}$

Image from Google Images

# Bayesian Learning

- Observations  $x$ , Parameters/Unobserved variables  $z$
- Likelihood =  $p(x | z)$

$$p\left(\text{ | \text{cat}\right)$$

Image from Google Images

- Observations  $x$ , Parameters/Unobserved variables  $z$
- Likelihood =  $p(x \mid z)$
- Prior =  $p(z)$

$p(\text{cat})$

# Bayesian Learning

- Observations  $x$ , Parameters/Unobserved variables  $z$
- Likelihood =  $p(x | z)$
- Prior =  $p(z)$
- Joint =  $p(x, z) = p(x | z)p(z)$



Image from Google Images

# Bayesian Learning

- Observations  $x$ , Parameters/Unobserved variables  $z$
- Likelihood =  $p(x | z)$
- Prior =  $p(z)$
- Joint =  $p(x, z) = p(x | z)p(z)$
- Marginal Likelihood =  $p(x) = \int_z p(x | z)p(z)dz$


$$p\left(\text{$$

Image from Google Images

# Bayesian Learning

- Observations  $x$ , Parameters/Unobserved variables  $z$
- Likelihood =  $p(x | z)$
- Prior =  $p(z)$
- Joint =  $p(x, z) = p(x | z)p(z)$
- Marginal Likelihood =  $p(x) = \int_z p(x | z)p(z)dz$
- Posterior =  $p(z | x)$

$$p(\text{cat} \mid \text{img})$$

Image from Google Images



# Bayesian Learning

- Observations  $x$ , Parameters/Unobserved variables  $z$
- Likelihood =  $p(x | z)$
- Prior =  $p(z)$
- Joint =  $p(x, z) = p(x | z)p(z)$
- Marginal Likelihood =  $p(x) = \int_z p(x | z)p(z)dz$
- Posterior =  $p(z | x)$
- Non-Bayesian Learning : Maximize likelihood wrt parameters  $z$

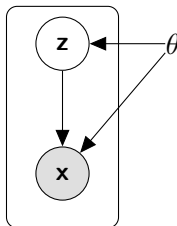
$$p(x | z)$$

- Bayesian Learning : Maximize marginal likelihood with a prior  $p(z)$  on parameters  $z$

$$p(x)$$

# Latent Variable Models

- Parameters:  $\theta$ , Latent variables:  $z$ , Observations  $x$
- Two step process:




- $z \sim p_{\theta}(z)$
- Given  $z$ ,  $x \sim p_{\theta}(x | z)$
- We only observe  $x \sim p_{\theta}(x)$
- $p_{\theta}(x) = \int_z p_{\theta}(x|z)p_{\theta}(z)dz.$

# Tasks

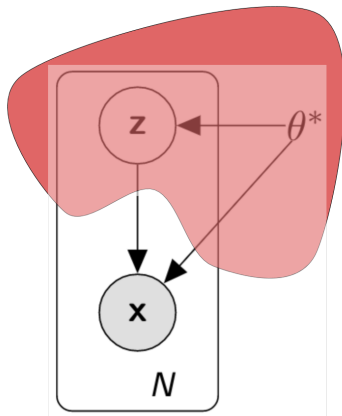
- ML/MAP inference on  $\theta$ .  
mimic data generation
- Posterior inference on  $z$  given  $x$ .  
coding/data representation.
- Marginal inference on  $x$   
denoising, inpainting, super-resolution.

Want efficient algorithms for all, with minimal assumptions.

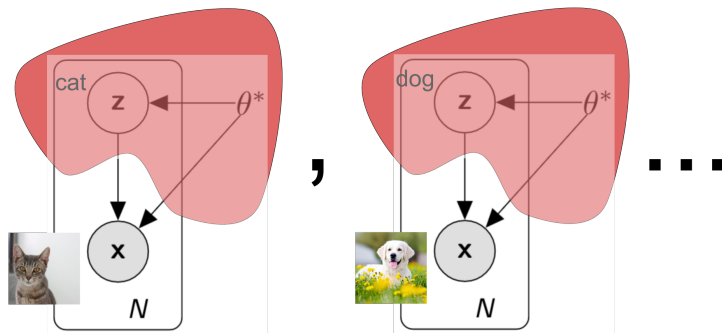


# Why is this hard?

- Lots of stuff hidden from us, we only see  $x$ .
- Data generation process can be complicated



# Ideas?



Given only images, we want

- $\theta$  close to  $\theta^*$
- An approximation to  $p_{\theta^*}(z | x)$
- A good model of  $p_{\theta^*}(x)$

Images from Google Images

# Some approaches

## Idea 1

- Integrate out  $z$ . Maximize marginal likelihood

$$p_{\theta}(\mathbf{x}) = \int_z p_{\theta}(\mathbf{x} | z) p_{\theta}(z) dz$$

.


# Some approaches

## Idea 1

- Integrate out  $z$ . Maximize marginal likelihood

$$p_{\theta}(\mathbf{x}) = \int_z p_{\theta}(\mathbf{x} | z) p_{\theta}(z) dz$$

Problem : nasty integral



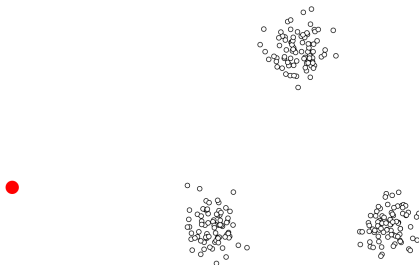
# Some approaches

## Idea 2

- Alternating optimization between  $z$  and  $\theta$  : Expectation Maximization.

Example : kMeans

- $\theta$  contains cluster centres.
- for every  $x_i$ ,  $z_i$  is the cluster id.





# Some approaches

## Idea 2

- Alternating optimization between  $z$  and  $\theta$  : Expectation Maximization.

Example : kMeans

- $\theta$  contains cluster centres.
- for every  $x_i$ ,  $z_i$  is the cluster id.

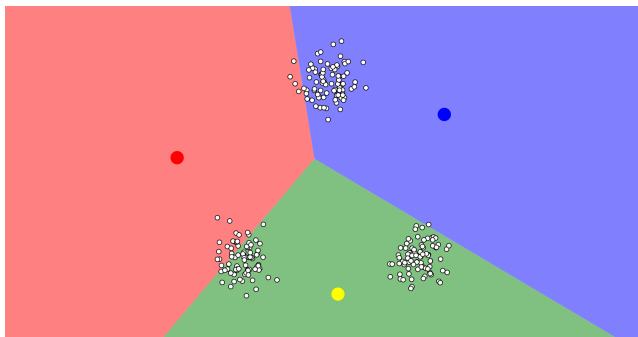


Image from <https://www.naftaliharris.com/blog/visualizing-k-means-clustering/>

# Some approaches

## Idea 2

- Alternating optimization between  $z$  and  $\theta$  : Expectation Maximization.

Example : kMeans

- $\theta$  contains cluster centres.
- for every  $x_i$ ,  $z_i$  is the cluster id.

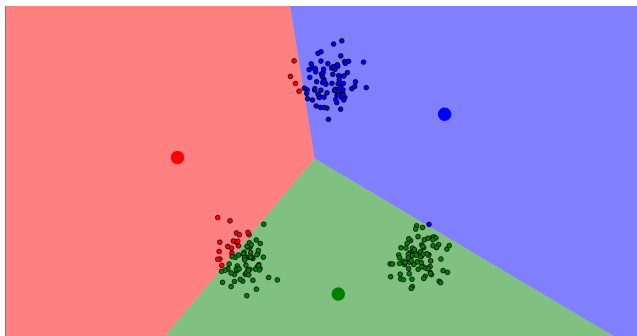


Image from <https://www.naftaliharris.com/blog/visualizing-k-means-clustering/>

# Some approaches

## Idea 2

- Alternating optimization between  $z$  and  $\theta$  : Expectation Maximization.

Example : kMeans

- $\theta$  contains cluster centres.
- for every  $x_i$ ,  $z_i$  is the cluster id.




Image from <https://www.naftaliharris.com/blog/visualizing-k-means-clustering/>

# Some approaches

## Idea 2

- Alternating optimization between  $z$  and  $\theta$  : Expectation Maximization.

Example : kMeans

- $\theta$  contains cluster centres.
- for every  $x_i$ ,  $z_i$  is the cluster id.

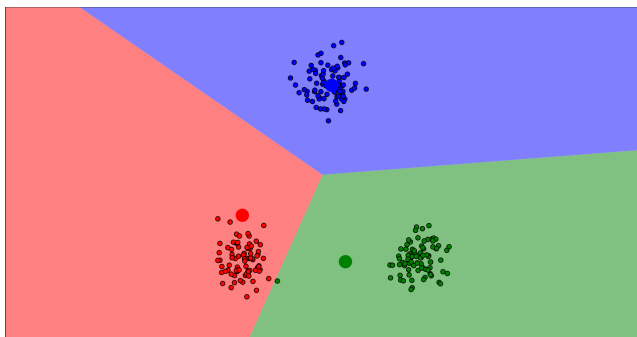


Image from <https://www.naftaliharris.com/blog/visualizing-k-means-clustering/>

# Some approaches

## Idea 2

- Alternating optimization between  $z$  and  $\theta$  : Expectation Maximization.

Example : kMeans

- $\theta$  contains cluster centres.
- for every  $x_i$ ,  $z_i$  is the cluster id.

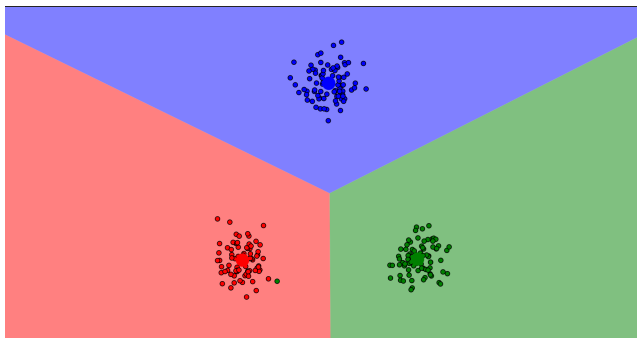


Image from <https://www.naftaliharris.com/blog/visualizing-k-means-clustering/>

# Some approaches

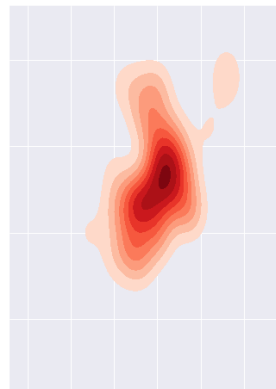
## Idea 2

- Alternating optimization between  $z$  and  $\theta$  : Expectation Maximization.

Problem : Posterior  $p_{\theta}(\mathbf{z} \mid \mathbf{x})$  may not be tractable.

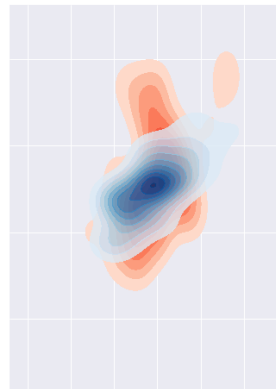
# Enter Variational Inference!

- Problem: We want to estimate some distribution  $p_{\theta}(\cdot)$ , but direct estimation is hard.



# Enter Variational Inference!


- Problem: We want to estimate some distribution  $p_{\theta}(\cdot)$ , but direct estimation is hard.
- Solution
  - Approximate  $p_{\theta}(\cdot)$  with a simpler distribution  $q_{\phi}(\cdot)$ .
  - Find parameters  $\phi$  such that the approximation is “close”.





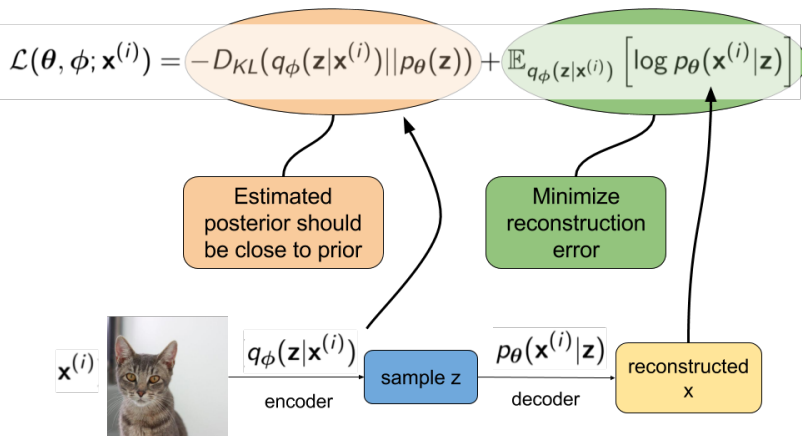
# Variational Inference – our setting

- Since the posterior  $p_{\theta}(\mathbf{z} \mid \mathbf{x})$  is intractable, use variational approximation  $q_{\phi}(\mathbf{z} \mid \mathbf{x})$ .
- $q_{\phi}(\mathbf{z} \mid \mathbf{x})$  : probabilistic **encoder**
- $p_{\theta}(\mathbf{x} \mid \mathbf{z})$  : probabilistic **decoder**



# Variational Lower Bound

- Using approximation leads to smaller marginal likelihood.
- Lower bound on marginal likelihood in terms of the variational approximation:



# Where is Deep Learning?

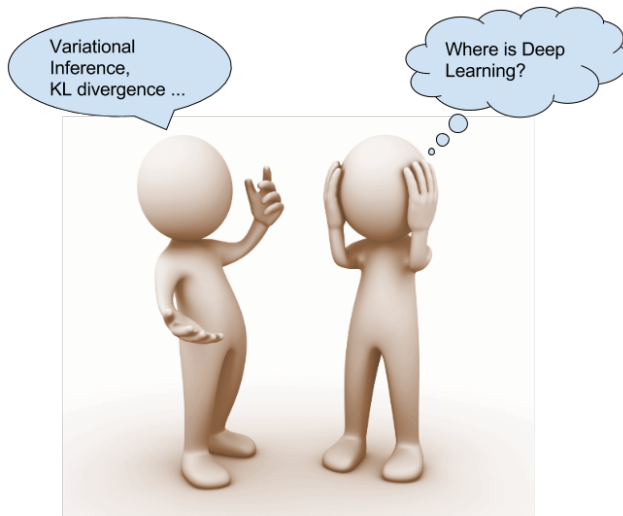


Image taken Google Images (modified)

# Parametrizing distributions

Ideas?

# Parametrizing distributions

## Ideas?

- Since we did not assume tractable posteriors, can use any arbitrary functions for generative and variational part.
- Only requirement - we should be able to optimize wrt  $\theta$  and  $\phi$ .
- For gradient based algorithms, we want a paramteric family which is differentiable wrt inputs and parameters.
- Can use neural networks.

# Parametrizing distributions

## Ideas?

- Since we did not assume tractable posteriors, can use any arbitrary functions for generative and variational part.
- Only requirement - we should be able to optimize wrt  $\theta$  and  $\phi$ .
- For gradient based algorithms, we want a parametric family which is differentiable wrt inputs and parameters.
- Can use neural networks.

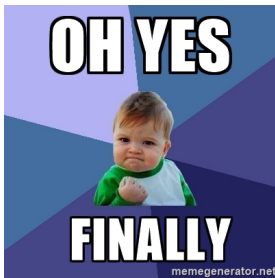
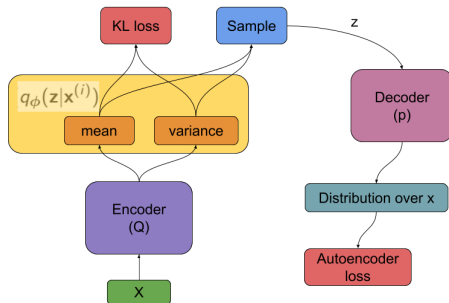


Image from Google Images

# Example : Variational Autoencoder

- $p_{\theta}(\mathbf{z}) = \mathcal{N}(\mathbf{z}; \mathbf{0}, \mathbf{I})$
- $p_{\theta}(\mathbf{x}|\mathbf{z})$  be a "simple" distribution whose distribution parameters are computed from  $\mathbf{z}$  with a neural network.
- Assume  $q_{\phi}(\mathbf{z}|\mathbf{x}^{(i)}) = \mathcal{N}(\mathbf{z}; \boldsymbol{\mu}^{(i)}, \boldsymbol{\sigma}^{2(i)}\mathbf{I})$ , where  $\boldsymbol{\mu}$ , and  $\boldsymbol{\sigma}$  are predicted using a neural network



# Gradient based optimization : Naïve method

- Lower bound is

$$\begin{aligned}\mathcal{L}(\theta, \phi; \mathbf{x}^{(i)}) &= -D_{KL}(q_\phi(\mathbf{z}|\mathbf{x}^{(i)})||p_\theta(\mathbf{z})) + \mathbb{E}_{q_\phi(\mathbf{z}|\mathbf{x}^{(i)})} \left[ \log p_\theta(\mathbf{x}^{(i)}|\mathbf{z}) \right] \\ &= \mathbb{E}_{q_\phi(\mathbf{z})} [f(\mathbf{z})]\end{aligned}$$

- Monte Carlo estimator:

$$\begin{aligned}\nabla_\phi \mathbb{E}_{q_\phi(\mathbf{z})} [f(\mathbf{z})] &= \mathbb{E}_{q_\phi(\mathbf{z})} \left[ f(\mathbf{z}) \nabla_{q_\phi(\mathbf{z})} \log q_\phi(\mathbf{z}) \right] \\ &\simeq \frac{1}{L} \sum_{l=1}^L f(\mathbf{z}^{(l)}) \nabla_{q_\phi(\mathbf{z}^{(l)})} \log q_\phi(\mathbf{z}^{(l)})\end{aligned}$$

$$\text{where } \mathbf{z}^{(l)} \sim q_\phi(\mathbf{z}|\mathbf{x}^{(i)})$$

- This has very high variance.



# Gradient based optimization : Naïve method

- Lower bound is

$$\begin{aligned}\mathcal{L}(\theta, \phi; \mathbf{x}^{(i)}) &= -D_{KL}(q_{\phi}(\mathbf{z}|\mathbf{x}^{(i)})||p_{\theta}(\mathbf{z})) + \mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x}^{(i)})} \left[ \log p_{\theta}(\mathbf{x}^{(i)}|\mathbf{z}) \right] \\ &= \mathbb{E}_{q_{\phi}(\mathbf{z})} [f(\mathbf{z})]\end{aligned}$$

- Monte Carlo estimator:

$$\begin{aligned}\nabla_{\phi} \mathbb{E}_{q_{\phi}(\mathbf{z})} [f(\mathbf{z})] &= \mathbb{E}_{q_{\phi}(\mathbf{z})} \left[ f(\mathbf{z}) \nabla_{q_{\phi}(\mathbf{z})} \log q_{\phi}(\mathbf{z}) \right] \\ &\simeq \frac{1}{L} \sum_{l=1}^L f(\mathbf{z}^{(l)}) \nabla_{q_{\phi}(\mathbf{z}^{(l)})} \log q_{\phi}(\mathbf{z}^{(l)}) \\ \text{where } \mathbf{z}^{(l)} &\sim q_{\phi}(\mathbf{z}|\mathbf{x}^{(i)})\end{aligned}$$

- This has very high variance. Q : Why?

# Gradient based optimization : Naïve method

- Lower bound is

$$\begin{aligned}\mathcal{L}(\theta, \phi; \mathbf{x}^{(i)}) &= -D_{KL}(q_{\phi}(\mathbf{z}|\mathbf{x}^{(i)})||p_{\theta}(\mathbf{z})) + \mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x}^{(i)})} \left[ \log p_{\theta}(\mathbf{x}^{(i)}|\mathbf{z}) \right] \\ &= \mathbb{E}_{q_{\phi}(\mathbf{z})} [f(\mathbf{z})]\end{aligned}$$

- Monte Carlo estimator:

$$\begin{aligned}\nabla_{\phi} \mathbb{E}_{q_{\phi}(\mathbf{z})} [f(\mathbf{z})] &= \mathbb{E}_{q_{\phi}(\mathbf{z})} \left[ f(\mathbf{z}) \nabla_{q_{\phi}(\mathbf{z})} \log q_{\phi}(\mathbf{z}) \right] \\ &\simeq \frac{1}{L} \sum_{l=1}^L f(\mathbf{z}^{(l)}) \nabla_{q_{\phi}(\mathbf{z}^{(l)})} \log q_{\phi}(\mathbf{z}^{(l)})\end{aligned}$$

where  $\mathbf{z}^{(l)} \sim q_{\phi}(\mathbf{z}|\mathbf{x}^{(i)})$

- This has very high variance. **Q : Why?**

A : Gradient of log of probability. Probability very close to zero means very large values.

# The reparametrization trick

- Problem in naïve method : learnable parameters responsible for producing probabilities.
- Observation: We don't need those probabilities, just an expectation taken using them.
- Solution
  - Instead of producing probability for each  $z$ , produce  $z$  directly
  - Make sure the distribution of  $z$  is the same.
  - For randomness in  $z$  generation, use a deterministic function with noise as input. i.e.

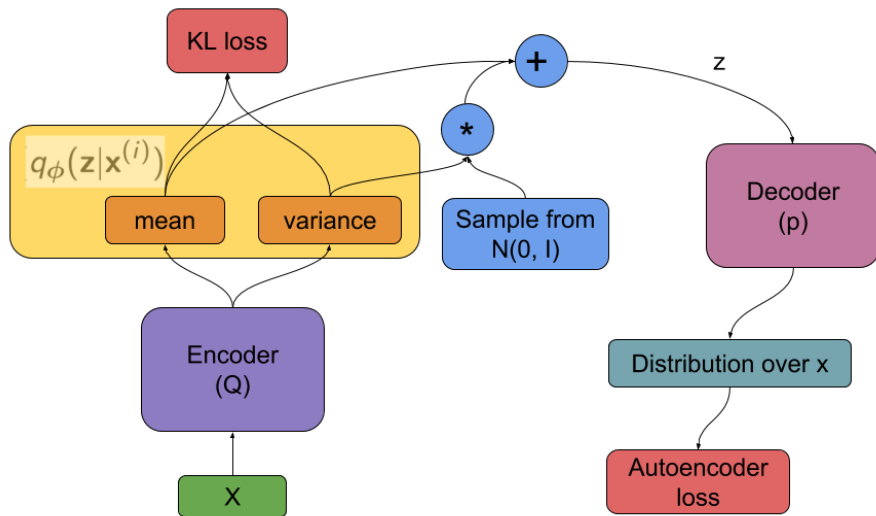
$$g_\phi(z, \epsilon) = \tilde{z} \sim q_\phi(z|x), \quad \epsilon \sim p(\epsilon)$$

- Now we can write expectations in terms of  $p(\epsilon)$ .

$$\mathbb{E}_{q_\phi(z|x^{(i)})} [f(z)] = \mathbb{E}_{p(\epsilon)} [f(g_\phi(\epsilon, x^{(i)}))] \simeq \frac{1}{L} \sum_{l=1}^L f(g_\phi(\epsilon^{(l)}, x^{(i)}))$$

$$\text{where } \epsilon^{(l)} \sim p(\epsilon)$$

# Reparametrization for VAE



# Putting it together

- Loss function:

$$\tilde{\mathcal{L}}(\theta, \phi; \mathbf{x}^{(i)}) = -D_{KL}(q_{\phi}(\mathbf{z}|\mathbf{x}^{(i)})||p_{\theta}(\mathbf{z})) + \frac{1}{L} \sum_{l=1}^L (\log p_{\theta}(\mathbf{x}^{(i)}|\mathbf{z}^{(i,l)}))$$

where  $\mathbf{z}^{(i,l)} = g_{\phi}(\epsilon^{(i,l)}, \mathbf{x}^{(i)})$  and  $\epsilon^{(l)} \sim p(\epsilon)$

- Algorithm

---

---

$\theta, \phi \leftarrow$  Initialize parameters

**repeat**

$\mathbf{X}^M \leftarrow$  Random minibatch of  $M$  datapoints (drawn from full dataset)

$\epsilon \leftarrow$  Random samples from noise distribution  $p(\epsilon)$

$\mathbf{g} \leftarrow \nabla_{\theta, \phi} \tilde{\mathcal{L}}(\theta, \phi; \mathbf{X}^M, \epsilon)$

$\theta, \phi \leftarrow$  Update parameters using gradients

**until** convergence of parameters  $(\theta, \phi)$

**return**  $\theta, \phi$

---

# Variational Autoencoder

## Experiments and Results

- Experiments on MNIST and FRAY face dataset
- Metric
  - Likelihood lower bound, same as optimization objective.

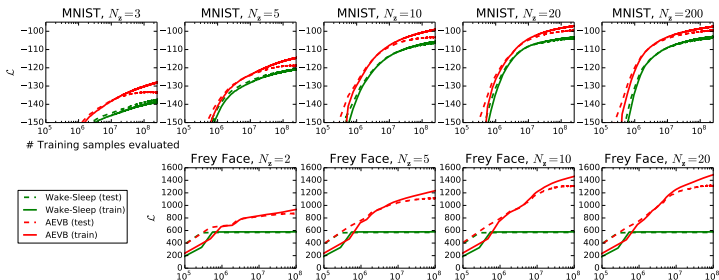


Image from <https://arxiv.org/abs/1312.6114>

# Variational Autoencoder

## Experiments and Results

- Experiments on MNIST and FRAY face dataset
- Metric
  - Estimated Marginal Likelihood

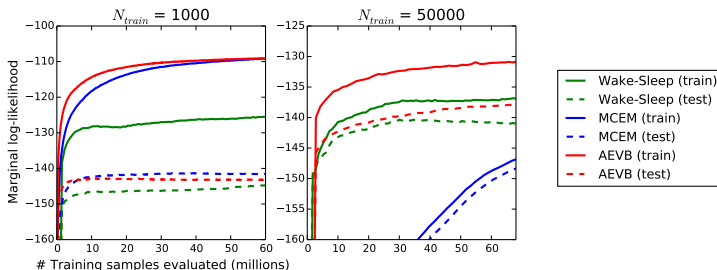
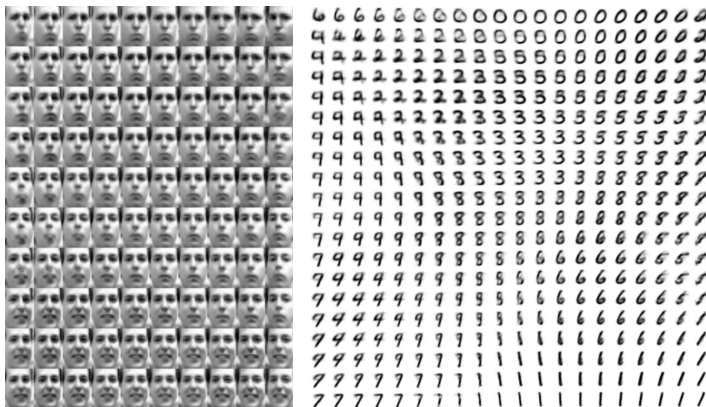


Image from <https://arxiv.org/abs/1312.6114>

# Variational Autoencoder

## Generated Samples



Images from <https://arxiv.org/abs/1312.6114>



# Variational Autoencoder

## Generated Samples

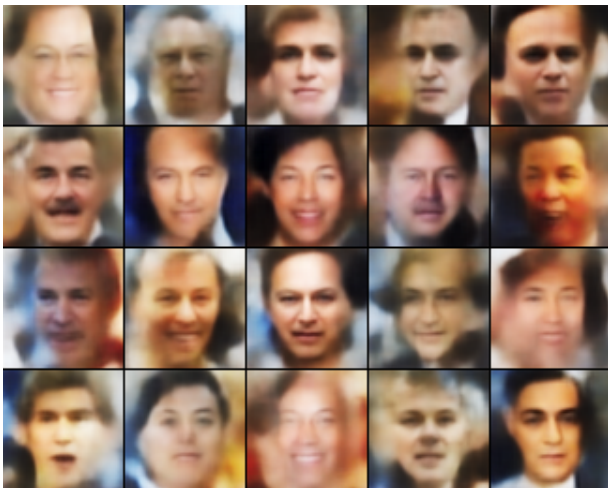


Image from <http://torch.ch/blog/2015/11/13/gan.html>

`https://transcranial.github.io/keras-js`

# Summary

- Problem : Latent variable model with intractable posterior, large dataset
- Solution
  - Variational approximation for posterior
  - Optimize variational lower bound
  - Reparametrize recognition model to reduce variance
  - Can plug in any function approximator (e.g. neural network)
- Example application
  - Variational Autoencoder

- What are the limitations of VAEs?
- How do they compare to GANs?
- VAE output images more blurred as compared to GANs
- Other questions?

# Thanks!